

Prisha Agarwal

925-964-5321 | prisha.agarwal@berkeley.edu | [linkedin.com/in/prishaa24](https://www.linkedin.com/in/prishaa24) | [prisha-agarwal.github.io](https://github.com/prisha-agarwal) | U.S. Citizen

EDUCATION

University of California, Berkeley

B.A. in Computer Science and Applied Mathematics

Berkeley, CA

Expected May 2028

Relevant Coursework: Data Structures & Algorithms, Artificial Intelligence, Computer Networks, Data Science, Abstract Linear Algebra, Abstract Algebra, Discrete Math, Object-Oriented Programming

TECHNICAL SKILLS

Languages & Frameworks: Python, JS, TS, Java; React, Redux, FastAPI, spaCy, XGBoost, MediaPipe, LangChain

Developer Tools: Git, Docker, AWS, GCP, PostgreSQL, Playwright, Vercel, OpenCV, ChromaDB, Supabase, Figma

EXPERIENCE

Snorkel AI

Software Engineer Intern

May 2026 – Aug. 2026

San Francisco, CA

- Shipped a **React/TypeScript** admin page that replaced scattered docs as the single editable source of metadata for 50+ current and all future projects, with **Redux-backed** per-section saves, **TanStack Router** unsaved-change guards, and duplicate-name checks for simultaneous edits; surfaced this to contracted experts in their own view
- Built the **Python/FastAPI** backend that resolves payment tiers for ~2,000 experts by country and entity, selects the currently active rate deterministically, and serves it via an access-controlled **REST API** endpoint
- Launched an in-app project update feed to replace manually sent out-of-band messages, with a **PostgreSQL** schema (SQLAlchemy + Alembic), a paginated endpoint, an admin posting page, and a searchable expert view

TurnPost AI

Machine Learning Engineer

Jan. 2026 – May 2026

Berkeley, CA

- Replaced GPT-4 API workflows with 5 deterministic NLP pipelines using **spaCy**, **TF-IDF**, and **Logistic Regression** for screenplay analysis, cutting runtime by 90% (30 → 3 min) and API cost to \$0 per script
- Built pipelines for character relationship graph extraction, emotion detection (90% accuracy on human-labeled test set), plot structure analysis, genre classification (80-85% accuracy across 18 categories), and entity recognition
- Fixed a critical character offset bug by replacing **str.index()** with **spaCy's ent.start_char**, correctly positioning repeated names, achieving 1,969 passing tests and fully deterministic and reproducible output

Krypton Consulting Inc.

Software Engineer Intern

Dec. 2025 – Jan. 2026

Newark, CA

- Built a secure authentication frontend with **React + Vite** and **Supabase Auth**, enabling gated access to an AI chat platform through email/password login, **OAuth** (Google / Facebook / Apple), and **password reset** flows
- Implemented persistent sessions via **Supabase JWT** storage with a **Remember Me** toggle that signs users out on next app load when unchecked, and deployed to **Vercel** for production with environment-variable config

San Jose State University

Machine Learning Researcher

Mar. 2024 – Aug. 2024

San Jose, CA

- Improved response relevance of LLM outputs by building a **Retrieval-Augmented Generation (RAG)** pipeline using **LangChain**, **OpenAI Embeddings**, and **ChromaDB** as the vector store, enabling **semantic retrieval** over domain-specific documents for accurate responses to real-world queries grounded in retrieved evidence
- Evaluated trustworthiness and interpretability of ML systems by conducting 20+ literature reviews/case studies on **Explainable AI (XAI)** studies across healthcare, finance, and law to identify deployment risks and best practices

PROJECTS

Humanoid Yoga Instructor | Finalist @ Physical AI Hackathon (Hosted by NVIDIA, Nebius, & More) Aug. 2026

- Built a real-time coaching pipeline that tracks 33 **MediaPipe** body landmarks from webcam video and computes joint angles to classify a user's specific pose alignment error, rather than just flagging "incorrect"
- Integrated motion-capture **SONIC** data with a **MuJoCo**-simulated **Unitree G1 humanoid** that physically demonstrates the targeted correction to guide the user to the correct pose in real time

Pebble | React, Python, ML, Computer Vision, VAPI (Voice Agent), MediaPipe, XGBoost, YOLOv8 Oct. 2025

- Built an intelligent study companion with a Pomodoro timer that tracks **behavioral signals** (including gaze, posture, hand movement, and phone usage) to detect loss of focus and deliver real-time voice reminders via **VAPI**
- Trained an **XGBoost** classifier on 4,000 self-collected labeled frames using **MediaPipe** facial landmarks, **OpenCV solvePnP** head pose estimation, and **YOLOv8** phone detection to output a real-time 0-100 focus score